

# Control-flow Hijacking Defences

**Instructor: Khaled Diab**

# Control-flow Hijacking Attacks so far...

---

- Buffer overflow: modify the return address
- Format string vulnerability: various range of attacks
- Heap overflows
- ...

# The Mistake!

---

## *Mixing code and data*

- Eventually, an attacker can inject code
- Source of other attacks...

# Defenses Overview

---

- Fix bugs
  - Automated tools
  - Rewrite software in different languages (examples?)
    - Legacy code?
- Run-time defenses:
  - StackGuard, Shadow Stack
- Platform defenses:
  - NOEXEC, ASLR

# StackGuard


---

- A technique that attempts to eliminate buffer overflow vulnerabilities
- A compiler modification
  - No source code changes
  - Requires recompiling the source code
- Patch for the function prologue and epilogue
- Prologue:
  - push an additional value into the stack (canary)
- Epilogue
  - pop the canary value from the stack and check that it hasn't changed



# Stack (no canary)

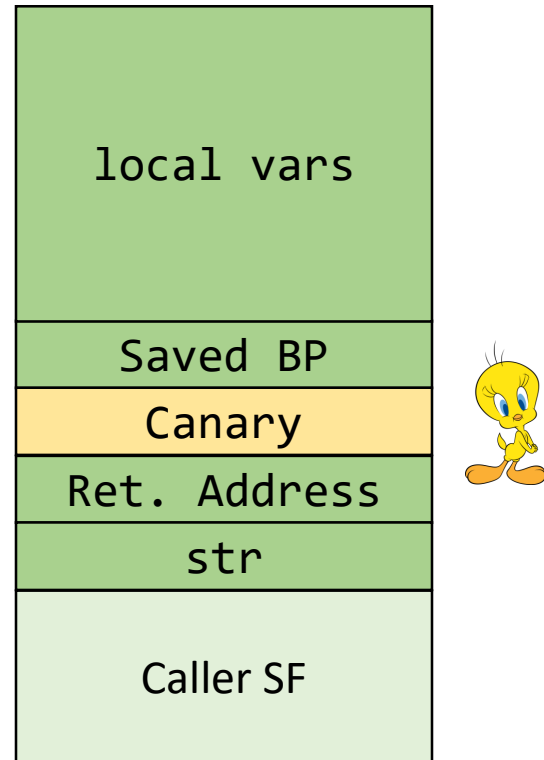
---



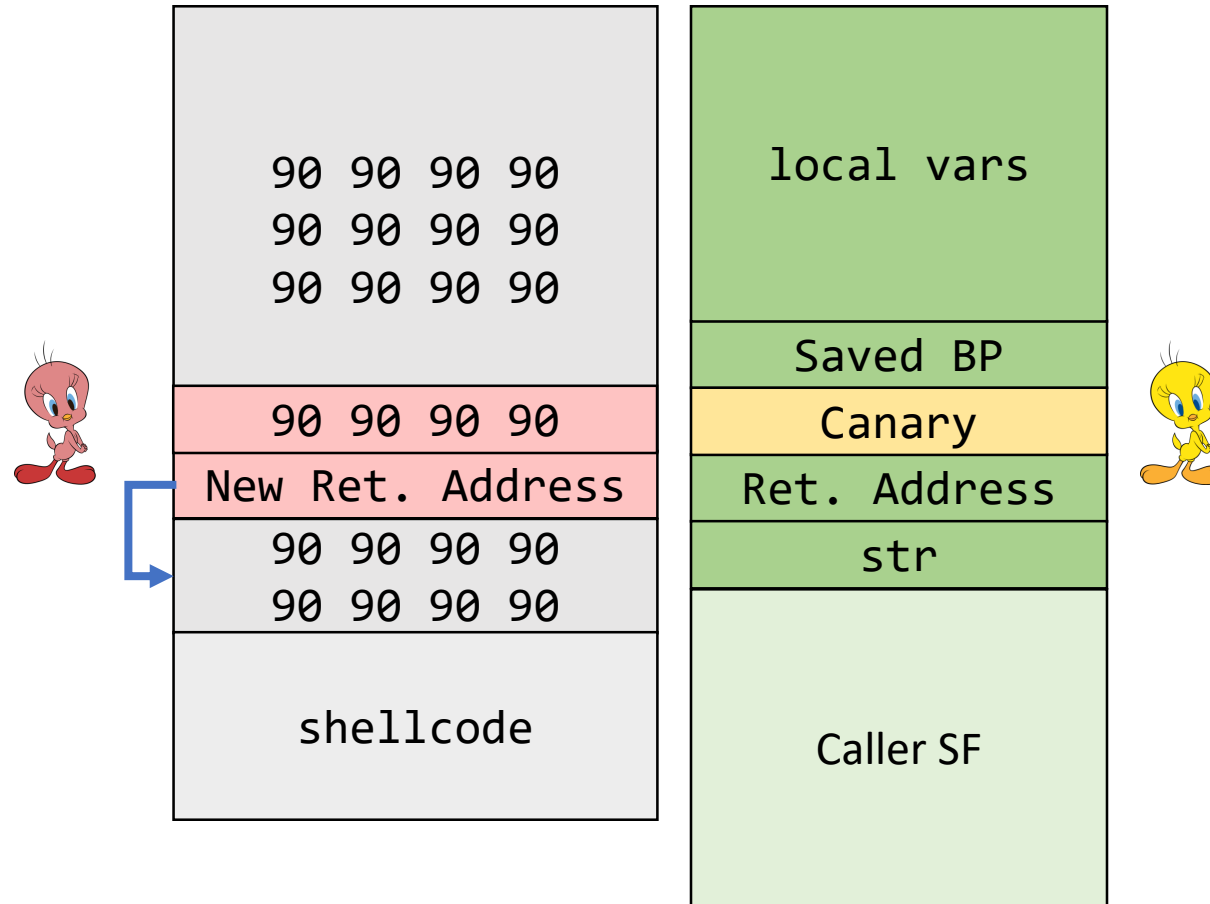
# Stack + Canary

---

Adds a random 32-bit value before the return address



# Stack + Canary (after overwriting ret. address)





# StackGuard Implementation in gcc

---

```
#include <stdio.h>

int main() {
    printf("Hello StackGuard");
    return 0;
}
```

```
$ gcc sg.c -o sg -fstack-protector-all
```

# StackGuard Implementation in gcc

---

```
0x0804846b <+0>:  lea    ecx,[esp+0x4]
0x0804846f <+4>:  and    esp,0xffffffff0
0x08048472 <+7>:  push   DWORD PTR [ecx-0x4]
0x08048475 <+10>: push   ebp
0x08048476 <+11>: mov    ebp,esp
0x08048478 <+13>: push   ecx
0x08048479 <+14>: sub    esp,0x14
0x0804847c <+17>: mov    eax,gs:0x14
0x08048482 <+23>: mov    DWORD PTR [ebp-0xc],eax
0x08048485 <+26>: xor    eax,eax
0x08048487 <+28>: sub    esp,0xc
0x0804848a <+31>: push   0x8048540
0x0804848f <+36>: call   0x8048330 <printf@plt>
```

# StackGuard Implementation in gcc

---

```
0x08048494 <+41>:    add     esp,0x10
0x08048497 <+44>:    mov     eax,0x0
0x0804849c <+49>:    mov     edx,DWORD PTR [ebp-0xc]
0x0804849f <+52>:    xor     edx,DWORD PTR gs:0x14
0x080484a6 <+59>:    je      0x80484ad <main+66>
0x080484a8 <+61>:    call    0x8048340
<__stack_chk_fail@plt>
0x080484ad <+66>:    mov     ecx,DWORD PTR [ebp-0x4]
0x080484b0 <+69>:    leave
0x080484b1 <+70>:    lea     esp,[ecx-0x4]
0x080484b4 <+73>:    ret
```

# Canary Types

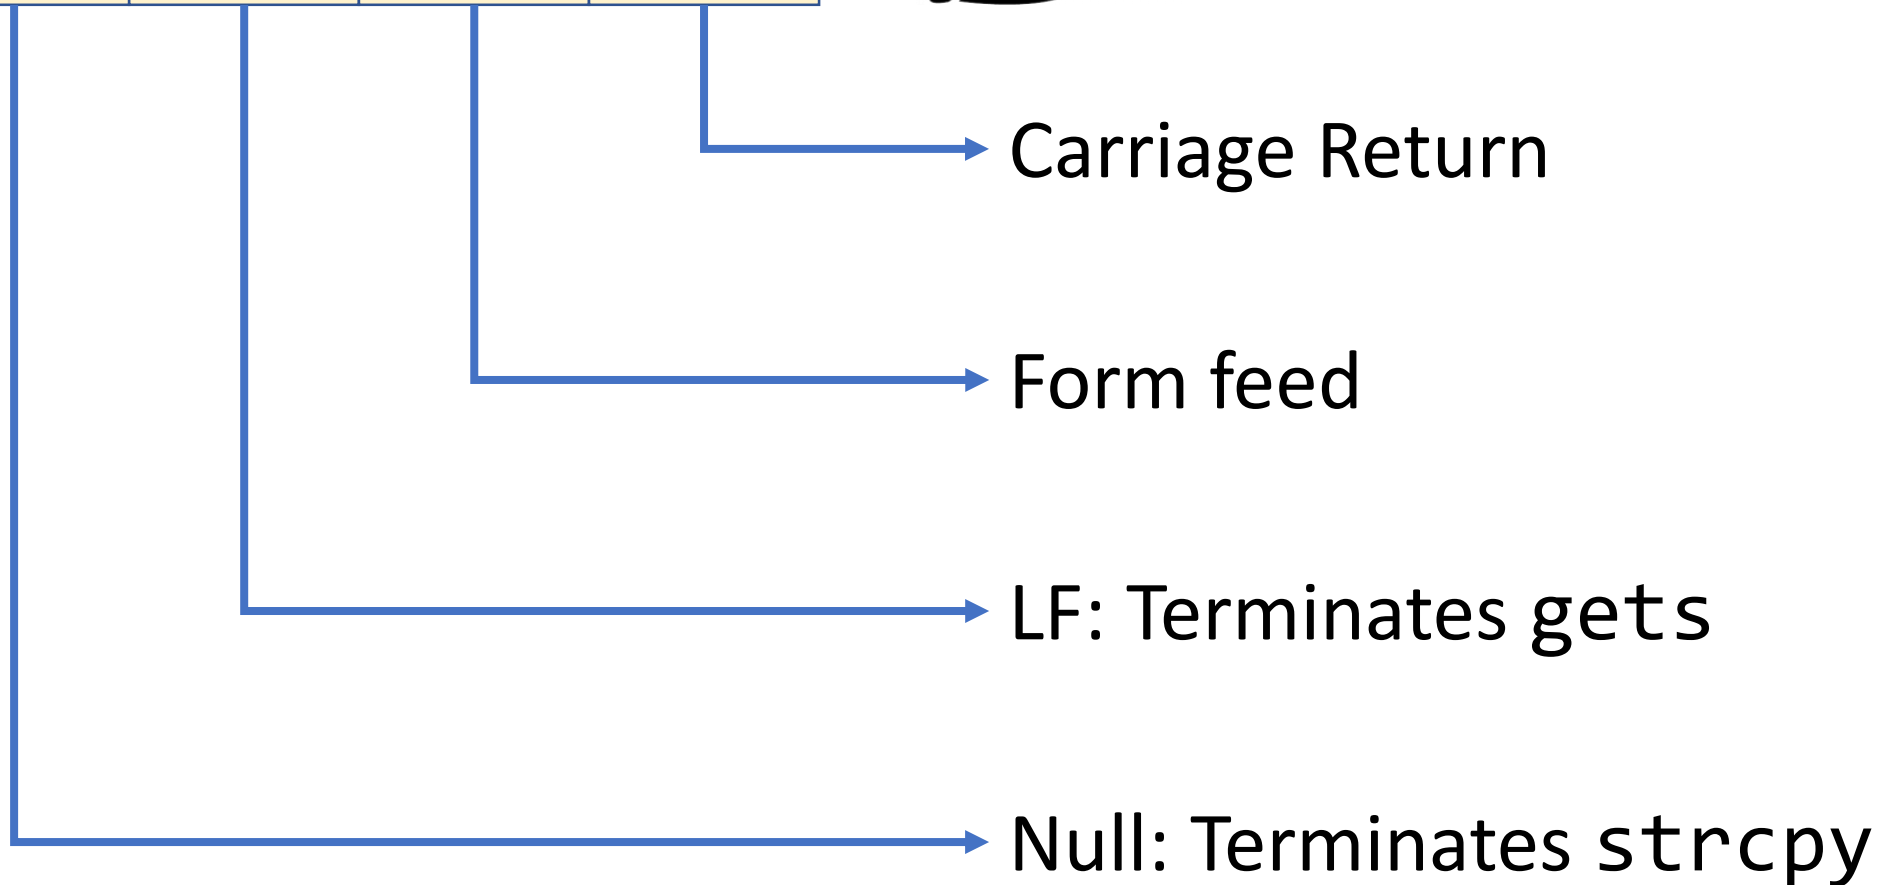

---

- Random Canary:
  - The original proposal
  - A 32-bit value
- Terminator Canary
  - A specific pattern
  - To act as string terminator for most string functions



# Terminator Canary

---



# Another Variation (Security vs Performance)

---

- gcc has two options:
  - -fstack-protector
    - Ignores some cases
  - -fstack-protector-all is very conservative
    - Adds protection to **all** functions
    - Performance overhead
- Chrome OS team has another proposal
  - -fstack-protector-strong
    - A superset of -fstack-protector
    - Examples: if a function has an array
    - [More details...](#)

# Shadow Stack

---

- Maintains return address at two stacks:
  - Original one: keeps SF information
  - Shadow: just the return address
- When a function returns, check

# Shadow Stack

## Traditional shadow stack

%gs:108

0xBEEF0048

Return address, R0  
Return address, R1  
Return address, R2  
Return address, R3

## Main stack

0x8000000

Parameters for R1  
Return address, R0  
First caller's EBP  
Parameters for R2  
Return address, R1  
EBP value for R1  
Local variables  
Parameters for R3  
Return address, R2  
EBP value for R2  
Local variables  
Return address, R3  
EBP value for R3  
Local variables

## Parallel shadow stack

0x9000000

Return address, R0  
Return address, R1  
Return address, R2  
Return address, R3



# What is the main assumption so far?

---


# What is the main assumption so far?

---

- The attacker can **only** overwrite the return address.
- Is that true?

# Stack-based Defenses: Limitations

- The attacker can modify local variables
  - Ones that are used in authentication
  - Function pointers
- The attacker can modify EBP
  - Frame pointer overwrite attack
  - EBP points to a fake frame inside the buffer
  - [More details](#)
- Assumes only the stack can be attacked!



# NOEXEC

---

- Only code segment executes code
- Set code segment to read-only
- Limitations:
  - Some applications need executable heaps
  - Can be bypassed using **Return-oriented Programming**
    - On Friday!

# Address Space Layout Randomization (ASLR)

```
$ sudo sysctl -w kernel.randomize_va_space=2
```

# Address Space Layout Randomization (ASLR)

---

- Map shared libraries to random location in process memory
  - Attacker cannot jump directly to execute function
- Consecutive runs result in different address space
- Need to randomize everything!
  - stack, heap, shared libs
- Discovering the address for shellcode becomes a difficult task
  - But not impossible!

# Address Space Layout Randomization (ASLR)

---

- Can be broken
- Heap Spray
  - The allocator is deterministic
  - If enough NOP+shellcode are sprayed in the heap, the attacker can make sure that the shellcode gets executed!

# Beyond Buffer Overflow Attacks

---

Consider this code:

```
int write(char* file, char* buffer) {  
    if (access(file, W_OK) != 0) {  
        exit(1);  
    }  
  
    int fd = open(file, O_WRONLY);  
    return write(fd, buffer, sizeof(buffer));  
}
```

- **Our goal:** open and write to regular file
- Code looks good!



# TOCTOU

---

- A race condition vulnerability

```
int write(char* file, char* buffer) {  
    if (access(file, W_OK) != 0) {  
        exit(1);  
    }  
    -----  
    int fd = open(file, O_WRONLY);  
    return write(fd, buffer, sizeof(buffer));  
}
```

An attacker can modify the file here! (how?)

In -sf /etc/passwd file  
00ps! What happened?

- The attacker now can modify a file they couldn't access before
- Recent incident: <https://duo.com/decipher/docker-bug-allows-root-access-to-host-file-system>

# Another Vulnerability

---

```
size_t len = readInt();  
char *buf;  
buf = malloc(len+9);  
read(fd, buf, len);
```

# Integer Overflow

---

```
size_t len = readInt();  
char *buf;  
buf = malloc(len+9);  
read(fd, buf, len);
```

What if len is large (e.g., 0xfffffffff)

→  $len+9 = 8$

→ The code allocates 8 bytes but can read a lot of data into buf

What if the variable controls access to a privileged operation?

# Another Vulnerability

---

```
char buf[80];  
void copyInput() {  
    int len = readInt();  
    char *input = readString();  
    if (len > sizeof(buf)) {  
        return;  
    }  
    memcpy(buf, input, len);  
}
```

```
void *memcpy(void *dst, const void * src, size_t n);
```

# Implicit Cast

---

Negative len can lead to large number of bytes being copied to buf!

```
char buf[80];  
void copyInput() {  
    int len = readInt();  
    char *input = readString();  
    if (len > sizeof(buf)) {  
        return;  
    }  
    memcpy(buf, input, len);  
}
```

```
void *memcpy(void *dst, const void * src, size_t n);
```

# How can we address these issues?

---

- Project ideas?

# Next lecture...

---

- return-to-libc
- ROP
- Control flow integrity